

End-to-End Machine Learning on Retail Transaction Data

A Theory-and-Results Report on the Sample Superstore Dataset

Debarshi Chakraborty

September 1, 2026

Companion codebase: `EY Training/ml_superstore/`
Data: `EY Training/data/Sample - Superstore.csv` (9,994 order lines, 2014–2017,
USA)

Abstract

This report documents an end-to-end machine learning study built on the Sample Superstore retail dataset. It follows the same five stages as the accompanying Jupyter notebooks and source package (`EY Training/ml_superstore/`): data ingestion and quality assessment, exploratory data analysis, feature engineering, time-series forecasting of monthly sales, and supervised learning for line-level profit prediction and loss detection. Unlike the notebooks, whose purpose is interactive exploration, this report is written to stand alone: every technique used is first introduced from first principles — the statistical or algorithmic theory it rests on — and only then applied to the data, with the resulting figures, tables and their business interpretation presented immediately afterwards. The headline findings are that (i) order discounting beyond roughly 30% is the dominant driver of loss-making order lines, far outweighing category, region or shipping mode; (ii) an additive Holt–Winters exponential smoothing model forecasts monthly sales with an out-of-sample MAPE of approximately 23%, competitive with a seasonal ARIMA model and clearly better than a seasonal-naïve baseline; and (iii) a random-forest classifier trained on leakage-safe features identifies loss-making order lines with a PR-AUC of 0.96 on a strictly future (2017) hold-out year, giving a concrete, auditable basis for a discount governance policy.

Contents

1	Introduction	1
1.1	Objective	1
1.2	Dataset	1
1.3	Software and reproducibility	2
1.4	Notation	2
2	Data Quality and Exploratory Analysis	3
2.1	Theory	3
2.1.1	Data quality auditing	3
2.1.2	Descriptive statistics and distribution shape	3
2.1.3	Correlation	3
2.1.4	Seasonal-trend decomposition (STL)	4
2.2	Method	4
2.3	Results	4
2.3.1	Data quality	4
2.3.2	Distributions	5
2.3.3	Profit by product hierarchy	6
2.3.4	Discount versus profit	7
2.3.5	Seasonality	9
2.3.6	Customers and products	9
2.4	Interpretation	10

List of Figures

2.1	Histograms (log-count y -axis) of Sales , Profit and Discount on the raw scale, showing pronounced right skew in Sales and a bimodal, discount-tier structure in Discount	5
2.2	Left: $\log(1 + \text{Sales})$ recovers an approximately unimodal, near-symmetric shape (Eq. (2.1) in the raw scale is large and positive; after the transform it is close to zero). Right: signed $\log(1 + \text{Profit})$, showing the loss-making mass as a distinct negative lobe rather than a long left tail.	5
2.3	Order-line counts by Ship Mode , Segment , Region and Category . Standard Class and the Consumer segment dominate; the four regions are comparably sized, with West largest.	6
2.4	Total profit by sub-category; red bars are net loss-makers.	7
2.5	Sales (left) and profit (right) cross-tabulated by Region and Segment. Sales are fairly even across cells; profit is not — the Central/Consumer cell is comparatively weak.	7
2.6	Left: line-level profit vs. discount (3,000-row sample), coloured by category — profit is scattered near zero at low discount and swings sharply negative as discount grows. Right: mean profit by discount band, crossing zero between the 10–20% and 20–30% bands and falling steeply beyond 30%.	7
2.7	Pearson correlation (Eq. (2.2)) among the core numeric fields. Sales–Profit : $\rho \approx 0.48$; Discount–Profit : $\rho \approx -0.22$ (negative but, per Fig. 2.6, understating the non-linear effect); Quantity is nearly uncorrelated with profit.	8
2.8	STL decomposition (Eq. (2.3), period 12) of monthly sales: a clear upward trend across the four years and a recurring within-year seasonal pattern that peaks late in the year, leaving a comparatively small, structure-free remainder.	9
2.9	Sales aggregated by month-of-year (left) and day-of-week (right, 0 = Monday). November–December is by far the strongest period; day-of-week shows only mild variation.	9

List of Tables

1.1	Raw schema of the Sample Superstore CSV.	2
2.1	Selected sub-categories by total profit (ascending). Full table has all 17 sub-categories.	6
2.2	Top 5 customers by total sales.	10

Chapter 1

Introduction

1.1 Objective

Retail order-line data — one row per product sold within an order — sits at the intersection of three classical machine-learning problem types: descriptive analytics (what happened), predictive analytics (what will happen), and diagnostic/prescriptive analytics (why it happened, and what to do about it). This report uses the public *Sample Superstore* dataset to walk through all three, mirroring a realistic analytics engagement:

1. **Data quality and exploratory analysis** — establish that the data is trustworthy and characterise its statistical structure.
2. **Feature engineering** — turn raw transactional fields into model-ready predictors, with explicit attention to information leakage.
3. **Forecasting** — model and project the aggregate monthly sales time series.
4. **Supervised learning** — model two line-level targets, *profit* (regression) and *is a line loss-making* (classification), and interpret what drives them.

Each stage in this report is structured identically: a **Theory** section derives or states the relevant statistical/ML machinery from first principles, followed by a **Method** section describing how it was applied to this dataset, then **Results** with the produced figures and tables, and finally an **Interpretation** section connecting the numbers back to a business reading. This mirrors, chapter for chapter, the five numbered notebooks in `EY_Training/ml_superstore/notebooks/` and the reusable modules in `src/`.

1.2 Dataset

The dataset contains $n = 9,994$ order lines spanning 4 years (2014-01-03 to 2017-12-30), 793 unique customers and 1,862 unique products, organised into 5,009 distinct orders across the United States. Table 1.1 summarises the raw schema.

Table 1.1: Raw schema of the Sample Superstore CSV.

Field	Type	Description
Row ID, Order ID	id	Line and order identifiers
Order Date, Ship Date	date	Order placement and shipment dates
Ship Mode	categorical	Standard / Second / First Class, Same
Customer ID, Customer Name	id/text	Buyer identity
Segment	categorical	Consumer / Corporate / Home Office
Country, City, State, Postal Code, Region	geography	U.S. location fields
Product ID, Category, Sub-Category, Product Name	catalogue	3-level product hierarchy
Sales	numeric (\$)	Line revenue
Quantity	numeric	Units sold
Discount	numeric [0,1)	Fractional discount applied
Profit	numeric (\$)	Line profit (can be negative)

1.3 Software and reproducibility

All analysis is implemented in Python 3.10 using `pandas`, `numpy`, `scikit-learn`, `statsmodels`, `matplotlib/seaborn`. Code lives in `EY_Training/ml_superstore/src/` as importable modules (`data_loader`, `features`, `forecasting`, `modeling`, `pipeline`), exercised interactively by six Jupyter notebooks and, for one-shot reproduction, by `python -m src.pipeline`. Every figure reproduced in this report was generated by that codebase and saved under `outputs/figures/`; every number quoted in the text is read directly from the corresponding CSV in `outputs/tables/` or the machine-written `outputs/REPORT.md`. No figure or statistic in this document was computed by hand.

1.4 Notation

Throughout, y_i denotes a target value for observation i , $\hat{y}_i$ its model prediction, $\mathbf{x}_i \in \mathbb{R}^p$ its feature vector, and n the sample size of the set under discussion (which may be a train, test, or full-data set, made explicit in context). For the monthly sales series we write y_t , $t = 1, \dots, T$, for the value in month t . $\mathbb{E}[\cdot]$, $\text{Var}(\cdot)$ and $\text{Cov}(\cdot, \cdot)$ denote expectation, variance and covariance under whatever probabilistic model is being assumed at that point (population model for theory, empirical/plug-in estimator when applied to data).

Chapter 2

Data Quality and Exploratory Analysis

2.1 Theory

2.1.1 Data quality auditing

Before any modelling, three questions must be answered about a tabular dataset $D = \{(\mathbf{x}_i)\}_{i=1}^n$: (i) is every field typed and parsed correctly, (ii) is any value missing, and (iii) is any row duplicated. Formally, for a column X we report the *missingness rate* $\hat{\pi}_X = \frac{1}{n} \sum_{i=1}^n \mathbb{1}[x_i \text{ is NA}]$ and *cardinality* $|\{x_i\}|$. A duplicate row under a primary key k is any pair $i \neq j$ with $k_i = k_j$; if such pairs exist, the key does not uniquely identify a record and downstream joins or grouped aggregates will silently double-count. These checks are cheap and guard against exactly the class of error (a re-exported CSV with duplicated header rows, an encoding mismatch producing NA in a text field) that otherwise corrupts every later stage silently.

2.1.2 Descriptive statistics and distribution shape

For a numeric column X with observed values $x_1, \dots, x_n$, the sample mean $\bar{x} = \frac{1}{n} \sum_i x_i$ and variance $s^2 = \frac{1}{n-1} \sum_i (x_i - \bar{x})^2$ summarise location and spread, but retail money variables are rarely well described by them alone because they are *right-skewed*: many small transactions and a long tail of large ones. Skewness is measured by the third standardised moment,

$$\gamma_1 = \frac{\frac{1}{n} \sum_i (x_i - \bar{x})^3}{\left(\frac{1}{n} \sum_i (x_i - \bar{x})^2\right)^{3/2}}, \quad (2.1)$$

with $\gamma_1 > 0$ indicating a right tail. A common remedy that stabilises both skew and the variance-vs-mean relationship (heteroscedasticity) seen in multiplicative processes like sales is the $\log(1+x)$ transform, used throughout this project on **Sales** (and, with a sign-preserving variant, on **Profit**) purely for *visualisation*; models are always trained on the original scale unless stated otherwise, since tree ensembles are invariant to monotone transforms and linear models' coefficients are more directly interpretable untransformed here.

2.1.3 Correlation

For two numeric variables X, Y , the Pearson correlation coefficient

$$\rho_{XY} = \frac{\text{Cov}(X, Y)}{\sigma_X \sigma_Y} = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_i (x_i - \bar{x})^2} \sqrt{\sum_i (y_i - \bar{y})^2}} \quad (2.2)$$

measures *linear* association, $\rho_{XY} \in [-1, 1]$. It is insensitive to non-linear relationships (e.g. a U-shape can have $\rho \approx 0$ despite strong dependence), which is precisely why Section 2.3.4 supplements the correlation matrix with a binned conditional-mean plot: discount and profit turn out to have a threshold-like, not linear, relationship.

2.1.4 Seasonal-trend decomposition (STL)

A time series y_t is classically modelled as the additive sum of a slowly-varying trend T_t , a periodic seasonal component S_t (period L , here $L = 12$ months), and a remainder R_t :

$$y_t = T_t + S_t + R_t. \quad (2.3)$$

STL (Seasonal-Trend decomposition using LOESS, Cleveland et al., 1990) estimates T_t and S_t by alternating two local regression (LOESS) smooths: it removes a running estimate of T_t from y_t to expose the seasonal cycle, smooths within each seasonal sub-series (all Januaries, all Februaries, ...) to obtain a smoothly time-varying S_t , then smooths $y_t - S_t$ to update T_t ; this is iterated to convergence. Unlike a single global seasonal index, STL lets the seasonal shape drift slowly over the four years of data, and unlike classical decomposition it is robust to outliers when the **robust** option down-weights large residuals between iterations. We use STL purely as a diagnostic here; the forecasting models of ?? are fitted directly on y_t , not on the decomposed components.

2.2 Method

The raw CSV is read with `src.data_loader.load_raw` (encoding auto-detection across UTF-8/Latin-1/CP1252) and typed with `clean`, which parses `Order Date`/`Ship Date`, strips whitespace from every text column, drops exact and Row ID duplicates, and derives `Ship Days` = `Ship Date` − `Order Date`. The quality audit, univariate/bivariate plots and STL decomposition are produced by notebooks `01_data_overview.ipynb` and `02_edu.ipynb`.

2.3 Results

2.3.1 Data quality

The audit (`outputs/tables/01_data_quality.csv`) finds **zero missing values in any of the 21 raw columns**, no fully duplicated rows, and no duplicate Row IDs after cleaning — the data is analysis-ready without imputation. `Order Date` spans 2014-01-03 to 2017-12-30 (1,237 distinct dates); `Ship Days` is always non-negative, i.e. no order ships before it is placed.

2.3.2 Distributions




Figure 2.1: Histograms (log-count y -axis) of **Sales**, **Profit** and **Discount** on the raw scale, showing pronounced right skew in **Sales** and a bimodal, discount-tier structure in **Discount**.




Figure 2.2: Left: $\log(1 + \text{Sales})$ recovers an approximately unimodal, near-symmetric shape (Eq. (2.1) in the raw scale is large and positive; after the transform it is close to zero). Right: $\text{signed } \log(1 + |\text{Profit}|)$, showing the loss-making mass as a distinct negative lobe rather than a long left tail.




Figure 2.3: Order-line counts by Ship Mode, Segment, Region and Category. Standard Class and the Consumer segment dominate; the four regions are comparably sized, with West largest.

2.3.3 Profit by product hierarchy

Table 2.1 (full detail in `outputs/tables/02_by_subcategory.csv`) aggregates profit by sub-category. **Tables** lose \$17,725 overall despite \$207k of sales, and **Bookcases** lose \$3,473 — both at average discounts above 21%, well above the store-wide mean. **Supplies** is also marginally loss-making. Every other sub-category is net profitable, led by Copiers (43% margin) and Phones (\$330k sales, 13% margin).

Table 2.1: Selected sub-categories by total profit (ascending). Full table has all 17 sub-categories.

Category	Sub-Category	Sales (\$)	Profit (\$)	Avg. Discount	Margin
Furniture	Tables	206,966	−17,725	26.1%	−8.6%
Furniture	Bookcases	114,880	−3,473	21.1%	−3.0%
Office Supplies	Supplies	46,674	−1,189	7.7%	−2.5%
... (14 profitable sub-categories) ...					
Technology	Phones	330,007	44,516	15.5%	13.5%
Technology	Copiers	149,528	55,618	16.2%	37.2%




Figure 2.4: Total profit by sub-category; red bars are net loss-makers.




Figure 2.5: Sales (left) and profit (right) cross-tabulated by Region and Segment. Sales are fairly even across cells; profit is not — the Central/Consumer cell is comparatively weak.

2.3.4 Discount versus profit




Figure 2.6: Left: line-level profit vs. discount (3,000-row sample), coloured by category — profit is scattered near zero at low discount and swings sharply negative as discount grows. Right: mean profit by discount band, crossing zero between the 10–20% and 20–30% bands and falling steeply beyond 30%.

Quantitatively, the loss rate ($\Pr(\text{Profit} < 0)$) is **3%** for undiscounted lines and **over 50%** for lines discounted 30% or more — an order-of-magnitude jump. Because the relationship is a threshold effect rather than linear, the raw Pearson correlation between **Discount** and **Profit** (visible in Fig. 2.7) understates its practical importance; this motivates using the categorical **discount_bucket** feature (rather than relying on a linear term alone) in ????




Figure 2.7: Pearson correlation (Eq. (2.2)) among the core numeric fields. **Sales–Profit:** $\rho \approx 0.48$; **Discount–Profit:** $\rho \approx -0.22$ (negative but, per Fig. 2.6, understating the non-linear effect); **Quantity** is nearly uncorrelated with profit.

2.3.5 Seasonality




Figure 2.8: STL decomposition (Eq. (2.3), period 12) of monthly sales: a clear upward trend across the four years and a recurring within-year seasonal pattern that peaks late in the year, leaving a comparatively small, structure-free remainder.




Figure 2.9: Sales aggregated by month-of-year (left) and day-of-week (right, 0 = Monday). November–December is by far the strongest period; day-of-week shows only mild variation.

2.3.6 Customers and products

The top-10 customers by sales (Table 2.2) generate between \$12k and \$25k each, but profitability is uneven within that list — *Sean Miller* is the single largest customer by revenue (\$25,043) yet is net loss-making (−\$1,981), while *Tamara Chand* generates comparable revenue at a much healthier margin. The same pattern repeats at product level: the *Cisco TelePresence System EX90* and two *GBC* binding systems are top-10 by revenue but net loss-making, again tracing back to discount depth on high-ticket Technology and Furniture items.

Table 2.2: Top 5 customers by total sales.

Customer	Sales (\$)	Profit (\$)	Orders
Sean Miller	25,043	−1,981	5
Tamara Chand	19,052	8,981	5
Raymond Buch	15,117	6,976	6
Tom Ashbrook	14,596	4,704	4
Adrian Barton	14,474	5,445	10

2.4 Interpretation

Three findings from this chapter shape every subsequent stage of the report. First, the data requires *no* missing-value or duplicate handling, so all modelling error downstream is genuinely about the signal, not data-cleaning artefacts. Second, **discount depth, not category or region, is the primary lever on profitability** — a threshold effect around 30% discount that is far stronger than any single-variable linear correlation suggests, which is why `discount_bucket` (??) and the corresponding permutation-importance results (??) recur as the dominant signal throughout the rest of the report. Third, the pronounced trend and December seasonality in Fig. 2.8 is exactly the structure a seasonal model must capture, and directly motivates the choice of Holt-Winters and SARIMA — both explicitly seasonal — over simpler alternatives in ??.